



# **Parallel & Distributed Computing**

## **Project Report**

**A parallel algorithm for constructing multiple independent spanning trees in bubble-sort networks**

### **Group Members:**

22i-1013 Mustafa Kamran

22i-0885 Malaika Afzal

22i-1074 Fatima Bakhsh

### **Section K**

# A parallel algorithm for constructing multiple independent spanning trees in bubble-sort networks

## Overview

This project investigates a parallel algorithm to construct multiple independent spanning trees (ISTs) in bubble-sort networks, evaluating its performance under different implementations:

- **Sequential (Python)**
- **MPI-based Distributed (Python + METIS)**
- **Hybrid MPI + OpenMP (C++)**

We focus on:

- Performance and scalability comparisons
- Bottleneck identification
- Optimization using METIS for partitioning and TAU for profiling

## 1. Algorithm Implementations

### 1.1 Sequential (Python)

- **Time Complexity:**  $O(n! \times n)$
- **Space Complexity:**  $O(n! \times n)$
- **Key Functions:**
  - `generate_bn(n)`
  - `parent1(v, t, n, inv, rpos)`
  - `build_trees(n)`

**Findings:**

- Suitable for  $n \leq 4$
- Performance degrades factorially with increasing  $n$
- Inefficient for larger graphs due to memory usage and computation time

## 1.2 MPI Implementation (Python + METIS)

- **Parallelism:** Distributed-memory model
- **Partitioning:** METIS used for graph-based load balancing
- **Key Functions:**
  - `write_metis_file()`
  - `run_metis()`
  - `optimize_partitioning()`

### Findings:

- Optimal performance when **partitions (k)  $\approx$  process count**
- Communication and inter-partition edges increase overhead
- Best for **medium-sized problems (n = 5–6)**

## 1.3 Hybrid MPI + OpenMP (C++)

- **Parallelism:** Combines distributed and shared-memory parallelism
- **Key Enhancements:**
  - OpenMP for intra-node parallelism
  - MPI for inter-node communication
  - TAU profiler used for identifying communication and cache bottlenecks

### Findings:

- Best performance for **large n** ( $\geq 7$ )
- Reduced communication overhead and better memory locality
- Performance scales well with threads and processes

## 2. Experimental Configuration

Test Type	Parameters
Sequential	n = 2 to 7
MPI	n = 2 to 7, processes = 2–4, k = 2–4
Hybrid	n = 5–7, processes = 1–4, threads = 2–6, k = 3–5

## 3. Performance Metrics

- **Execution Time**
- **Speedup:**  $\text{Speedup} = T_1 / T_p$
- **Efficiency:**  $\text{Efficiency} = \text{Speedup} / p$
- **Partitioning Efficiency**
- **Cache Metrics (L1, L2, L3)**
- **Instruction & Cycle Counts**
- **Call Path Profiling (TAU)**

## 4. Scalability Analysis

### 4.1 Strong Scaling

- Measures improvement with more resources at constant problem size
- Efficiency drops with excess parallelism due to communication overhead

## 4.2 Weak Scaling

- Measures efficiency as problem size increases with resources
- Hybrid model shows better sustained performance

## 5. Implementation Comparison

Feature	Sequential	MPI	Hybrid (MPI + OpenMP)
Best Use Case	$n \leq 3$	$n = 4-6$	$n \geq 7$
Parallelism	None	Distributed	Hybrid
Scalability	$\times$	Moderate	High
Memory Usage	Very High	Moderate	Efficient
Communication Overhead	None	High	Low
Code Complexity	Low	Medium	High

## 6. Detailed Performance Analysis

### 6.1 Small Problems ( $n < 4$ )

- Sequential performs best
- Parallel overhead dominates in MPI and Hybrid

### 6.2 Medium Problems ( $n = 4-6$ )

- MPI shows 2x–4x speedup
- Hybrid reaches up to 6x speedup due to thread-level optimization

### 6.3 Large Problems ( $n \geq 7$ )

- Sequential is impractical

- MPI is memory-intensive
- Hybrid offers highest throughput and scalability

## 7. Optimization Insights

### 7.1 Partitioning

- METIS enhances load balancing
- Ideal when  $k = \text{number of MPI processes}$
- Fewer inter-partition edges yield better performance

### 7.2 Threading & Memory

- OpenMP improves cache locality and reduces memory traffic
- Use 2–4 threads for medium  $n$ , 4–6 for large  $n$
- Shared-memory use improves efficiency

### 7.3 Profiling (TAU)

- Identified hotspots in MPI communication
- Cache performance improves significantly with hybrid model
- Revealed optimal process/thread combinations

## 8. Recommendations

Problem Size	Recommended Implementation	Notes
$n \leq 3$	Sequential	Minimal overhead
$4 \leq n \leq 6$	MPI	Ensure good partitioning
$n \geq 7$	Hybrid (MPI + OpenMP)	Best resource utilization

### Resource Configuration Tips:

- Match **MPI processes to partitions**
- Use **2–6 OpenMP threads per process**
- Optimize based on cache and core count

## 9. Conclusion

The **Hybrid MPI + OpenMP implementation** demonstrates superior scalability, memory efficiency, and runtime performance, especially for large-scale problems. The combination of METIS-based partitioning and TAU-guided tuning enables significant optimization across both computation and communication layers.

The correct choice of implementation depends on:

- Problem size
- Hardware architecture
- Available memory and cores

By carefully selecting configurations and tools, our algorithm can scale effectively from small systems to high-performance clusters.